

```

int cur_buf = 0; // Selects which buffer we are currently using.

do {
    unsigned int offset_for_next_entity =
        address_of_next_entity_to_process - start_of_nodel_heap;

    /* Read an entity from Node 1's memory, via software cache.
       We use a blocking operation because we need the result to
       continue processing, and we hope the cache will mean that the
       result is held locally and will thus arrive quickly. */
    mrapi_rmem_read(sw_cache_hndl,
                    offset_for_next_entity,
                    &next_entity_to_process,
                    0,
                    sizeof(Entity),
                    1, /* num_strides is 1 */
                    0, /* rmem_stride is irrelevant */
                    0, /* local_stride is irrelevant */
                    &status);

    // CHECK STATUS FOR ERROR
    if (status != MRAPI_SUCCESS) {
        ERR("Unable to read remote memory sw cache");
    }

    result_buffers[cur_buf][num_entities_processed % BUFFER_SIZE] =
        process( & next_entity_to_process );
    num_entities_processed++;

    address_of_next_entity_to_process =
        (unsigned int)(next_entity_to_process.next);

    if((num_entities_processed % BUFFER_SIZE) == 0)
    {

        // CHECK STATUS FOR ERROR - DETAILS OMITTED

        /* Issue non-blocking DMA of buffer-full of results back
           to Node 1's memory. We use a non-blocking operation
           because we do not need to wait for the write to
           complete in order to continue processing the list: it is
           preferable to overlap communication with computation. */
        mrapi_rmem_write_i(
            dma_hndl,
            num_entities_processed*sizeof(float),
            result_buffers[cur_buf],
            0,
            BUFFER_SIZE*sizeof(float),
            1, /* num_strides is 1 */
            0, /* rmem_stride is irrelevant */
            0, /* local_stride is irrelevant */

```